

Algorithms and Data Structures



COMP261

Pathfinding 1: Dijkstra's Algorithm

Yi Mei

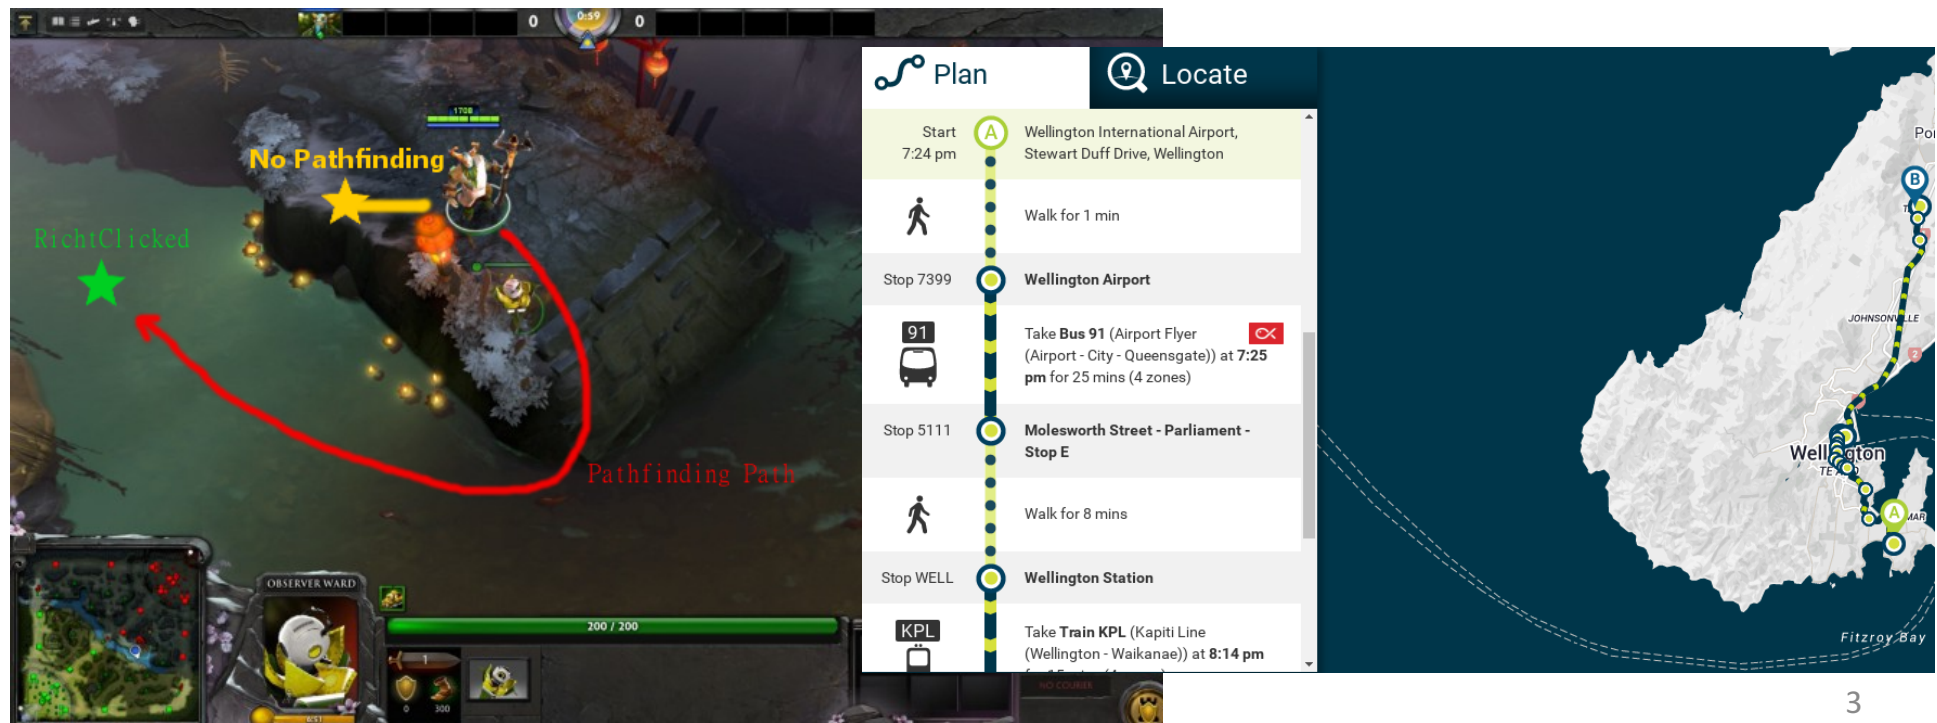
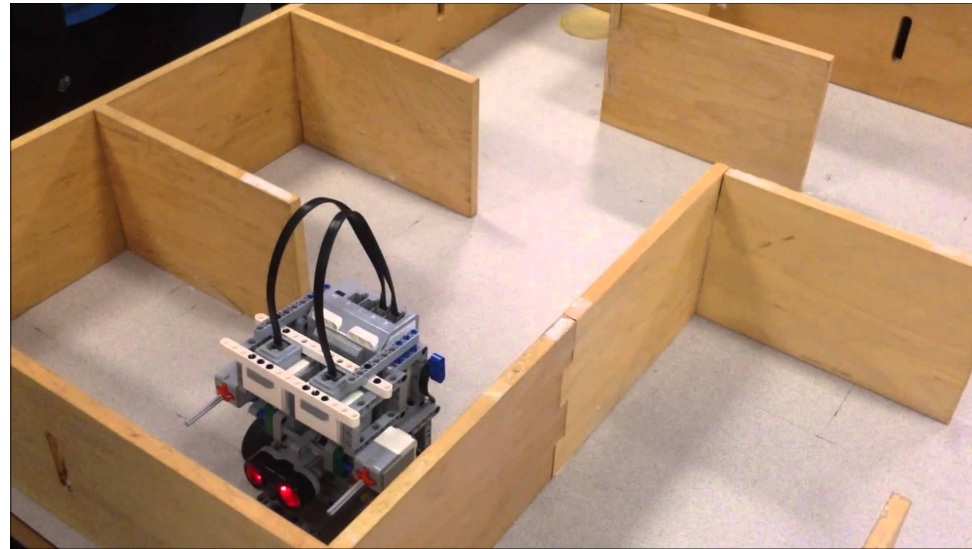
yi.mei@ecs.vuw.ac.nz

Outline

- Path finding
- Graph search technique for path finding
- Dijkstra's algorithm

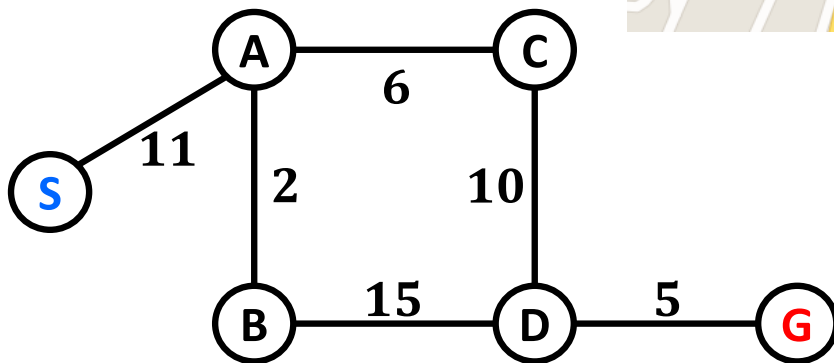
Path Finding

- Robotics
- Video games
- Journey planner



Path Finding

- Model environment as **graph**
- From VUW to Courtney PI?

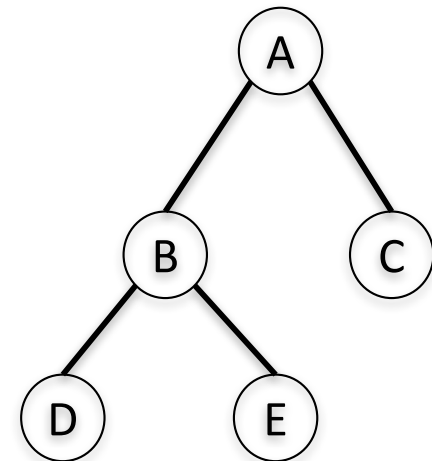


Path Finding

- Given a graph with a set of nodes and a set of edges, each edge is **undirected** and has a **cost** (e.g. length/travel time)
- Find the **least-cost** (e.g. shortest/fastest) path(s)
 - From **one node to another** (1-to-1)
 - From **one node to all other nodes** (1-to-all)
 - For **all node pairs** (all-to-all)
- Graph search
 - Depth-First Search (DFS)
 - Breadth-First Search (BFS)
 - ...

Graph Search

- Keep a “**fringe**” (all the leaf nodes of the search tree: frontier)
 - **Data structure** for the fringe: stack/queue/priority queue
- Start the fringe with one node (root)
- **Repeat until stopping criteria** are met:
 - **Choose** a node from the fringe to visit (visit a leaf node);
 - **Add** its neighbours to the fringe, and remove the node from fringe (expand the leaf node and grow the search tree);
- Example: DFS to visit all nodes
 - Use **stack** as fringe (last-in-first-out)
 - Step 1: fringe = {A}, visited = {}
 - Step 2: visit A; fringe = {B,C}, visited = {A}
 - Step 3: visit C; fringe = {B}, visited = {A,C}
 - Step 4: visit B; fringe = {D,E}, visited = {A,C,B}
 - Step 5: visit E; fringe = {D}, visited = {A,C,B,E}
 - Step 6: visit D; fringe = {}, visited = {A,C,B,E,D}
- What data structure for fringe for BFS?



Dijkstra's Algorithm

- 1-to-all path finding (from one node to all other nodes)
- Also for 1-to-1 (by early stopping)
- Use the graph search technique
 - Start the fringe with the start node (`start`)
 - Always `expand` the `unvisited` node that are “most likely to be in the shortest path” first
 - It has the **minimal cost from the start node**
 - **Expand**: visit, remove from the fringe, and add its neighbours into the fringe
- For this purpose, a data structure `SearchNode`, containing
 - `Node`: the node in the graph it represents
 - `Cost`: the minimal cost from the start node found so far
 - `Prev`: the previous node in the shortest path found so far
 - Denoted as `<Cost, Node, Prev>`

Dijkstra's Algorithm

Input: A weighted graph and a *start* node

Output: Shortest paths from *start* to all other nodes

Initially all the nodes are *unvisited*, and the fringe has a single element *<0, start, null>*;

While (*the fringe is not empty*) {

 Expand *<cost*, node*, prev*>* from the fringe, where **cost* is the minimal cost** among all the elements in the fringe;

if (*node* is unvisited*) {

 Set node* as *visited*, and set *node*.prev = prev**;

for (*edge = (node*, neigh)* outgoing from node*) {

if (*neigh is unvisited*) {

costToNeigh = cost + edge.weight*;

 add a new element *<costToNeigh, neigh, node*>* into the fringe;

 }

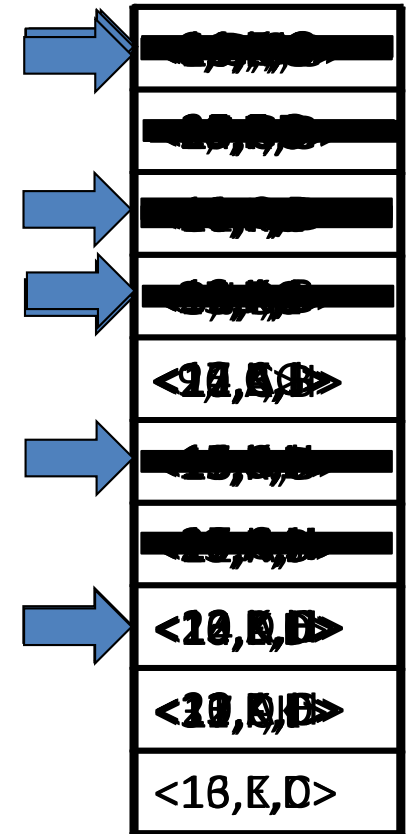
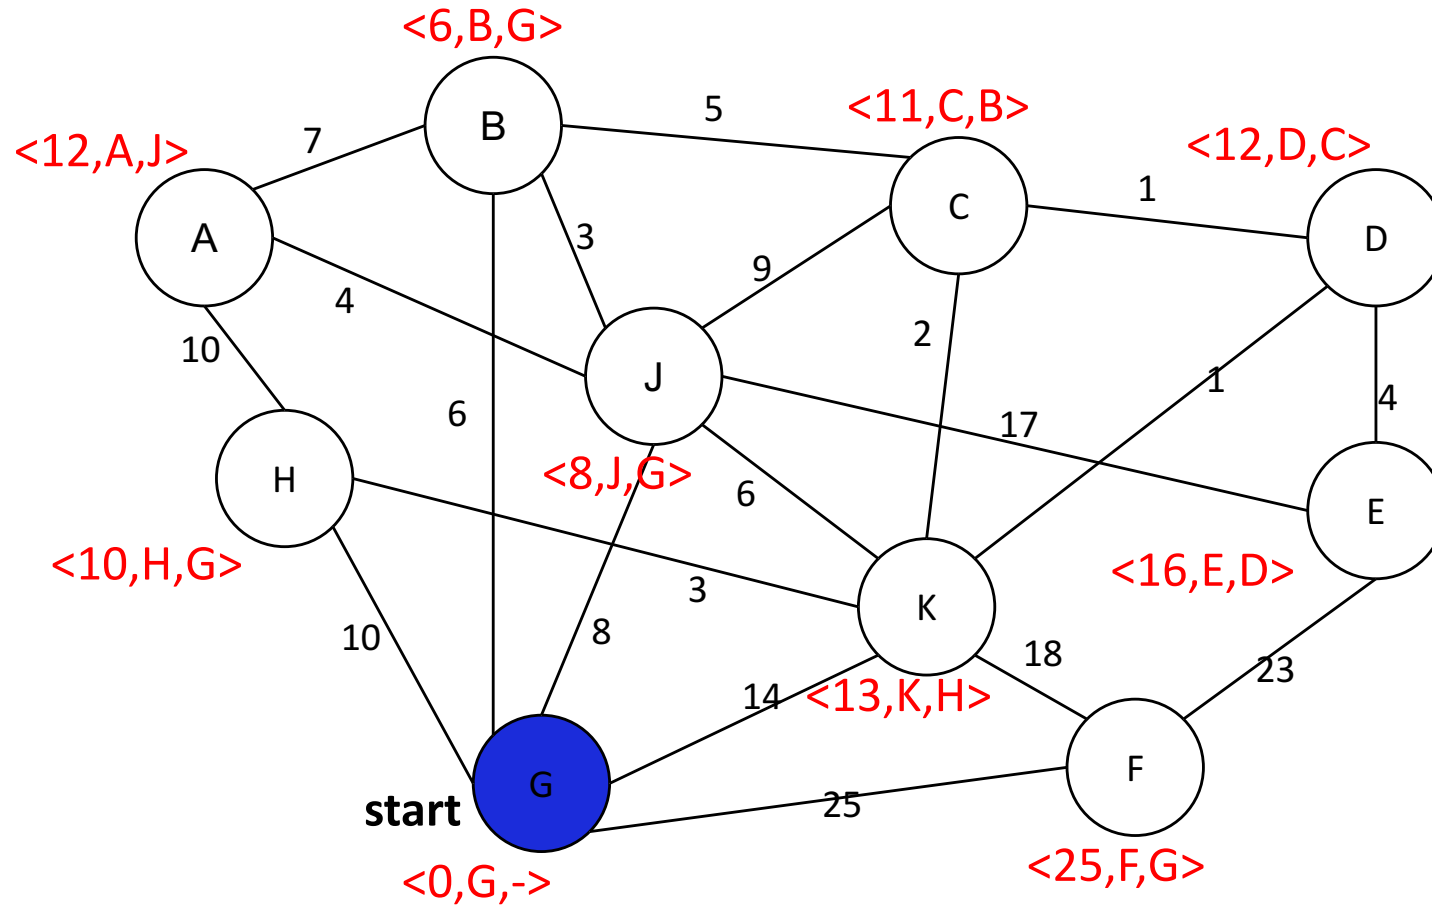
 }

 }

}

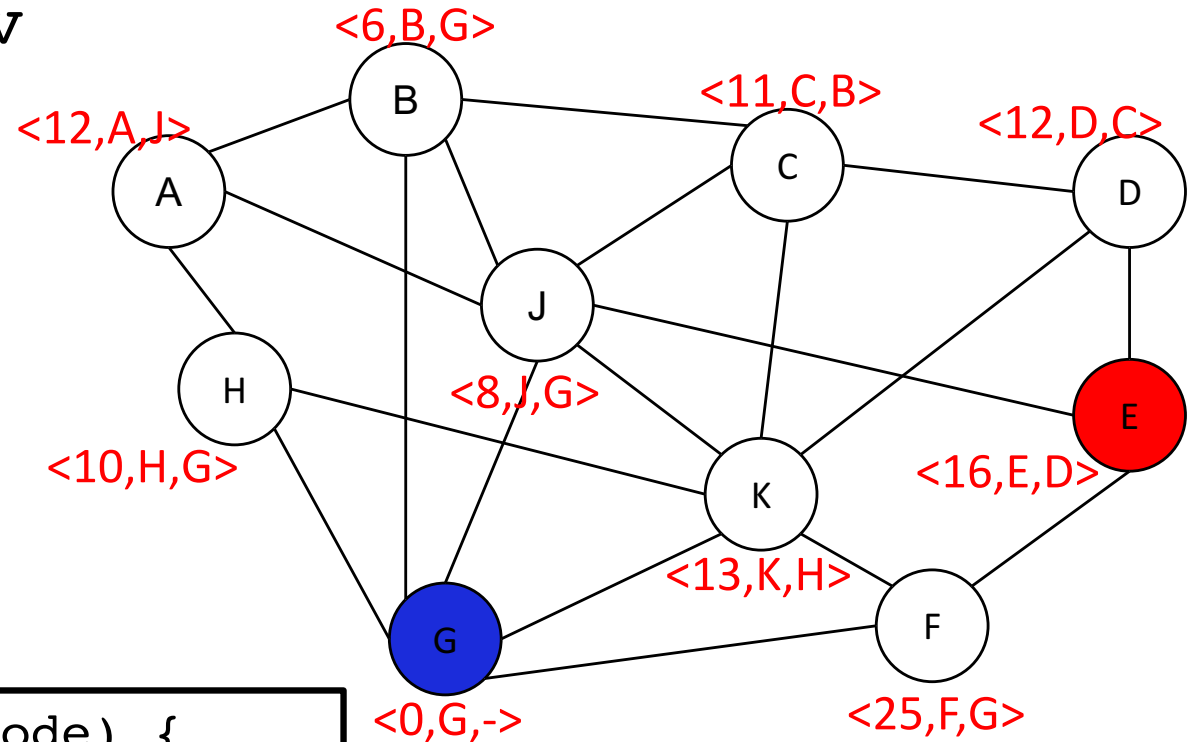
Obtain the shortest path based on the *.prev* fields;

Example of Dijkstra's Algorithm



Obtain Shortest Path

- We can obtain the shortest path from the start node to any other node from prev



```
getShortestPath(start, node) {  
    path ← (node), curr ← node;  
    while (not curr.prev = start) {  
        curr ← curr.prev;  
        path ← (curr, path);  
    }  
    return path;  
}
```

Path

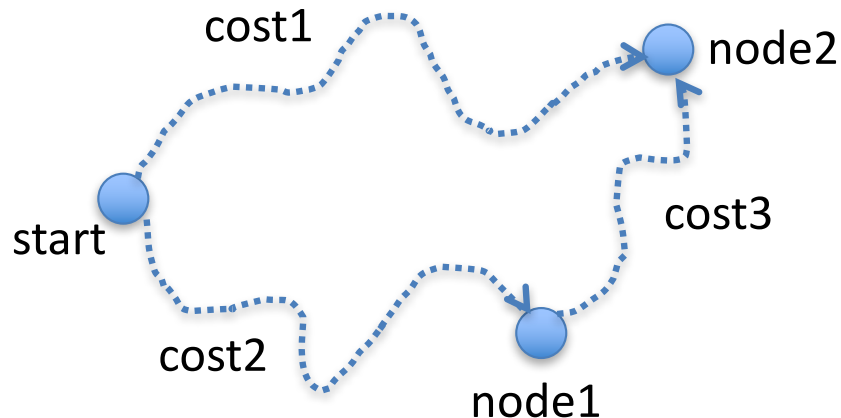


Correctness of Dijkstra's Algorithm

- **Theorem:** the path found for each node by Dijkstra's algorithm is the shortest path from the start node to the node
- **Proof:**
 - Each node is **visited only once**, when it is removed from the fringe
 - If future SearchNode contains the same node, it will be skipped
 - The **minimal cost** from the start node to a node is obtained when **visiting the node (for the first time)**

Correctness of Dijkstra's Algorithm

- **Theorem:** the path found for each node by Dijkstra's algorithm is the shortest path from the start node to the node
- **Proof (continued):**
 - No shorter path from the start node to a node **before** it is visited.
 - Otherwise the node would have been visited earlier (Dijkstra's algorithm always expand a node that is nearer the start node)
 - No shorter path from the start node to a node **after** it is visited.
 - Dijkstra's algorithm always expand a node that is nearer the start node



If $\text{cost1} < \text{cost2} + \text{cost3}$, then
 $\langle \text{cost1}, \text{node2}, \text{start} \rangle$ is visited **before**
 $\langle \text{cost2} + \text{cost3}, \text{node2}, \text{node1} \rangle$

If $\text{cost1} > \text{cost2} + \text{cost3}$, then
 $\langle \text{cost1}, \text{node2}, \text{start} \rangle$ is visited **after**
 $\langle \text{cost2} + \text{cost3}, \text{node2}, \text{node1} \rangle$

Data Structure for Fringe

- Which **data structure for the fringe** should be used in Dijkstra's Algorithm?
- Most common operation to the fringe
 - Add an element **<cost, node, prev>** to the fringe
 - Visit the node **<cost*, node*, prev*>** with the **minimal cost among all the elements** in the fringe (and remove it from the fringe)
- **Priority queue** – treat cost as the priority
 - Efficient for getting the element with the best priority

1-to-1 Dijkstra's Algorithm

- If we want to find **ONLY** the shortest (least cost) path from the **start node to a particular goal node**, then we can do it **faster**
 - Stop once we find the shortest path to the goal node, no need to continue for all the other nodes

Input: A weighted graph, a **start** node, a **goal** node

Output: A shortest path from **start** to **goal**

Initially all the nodes are **unvisited**, and the fringe has a single element **<0, start, null>**;

While (*the fringe is not empty*) {

 Expand **<cost*, node*, prev*>** from the fringe;

if (*node* is unvisited*) {

 Set node* as **visited**, and set **node*.prev = prev***;

if (node* is the goal node) **return**;

 // add all the unvisited neighbours of node* into the fringe

 ...

 }

}

Summary

- Path finding
 - 1-to-1 (Early-stop Dijkstra's algorithm)
 - 1-to-all (Dijkstra's algorithm)
 - All-to-all (not mentioned, but can also be solved by Dijkstra's algorithm)
- Dijkstra's algorithm
 - Graph search (fringe to store leaf nodes of search tree)
 - Priority queue for fringe
 - Always visit the node with minimal cost from the start node
- Can we do better?
 - Next lecture: A* search