



# COMP261 Lecture 24

## B and B+ Trees

**Victoria**  
UNIVERSITY OF WELLINGTON

*Te Whare Wānanga  
o te Ūpoko o te Ika a Māui*




CAPITAL CITY UNIVERSITY

# B Trees are great – can we do better?

- B-Trees are great when they contain a set of values
  - Only need to store the values you are searching on.
- In file/database setting:
  - Need to store key-value pairs
  - Search down the tree governed only by the keys
  - Values usually much larger than keys (eg, a whole record from a database table)
  - Storing values with keys reduces number of keys in each node (assuming fixed size nodes).
    - ⇒ lower branching factor
    - ⇒ deeper trees
    - ⇒ slower access times
- What if we want to index on more than one key?

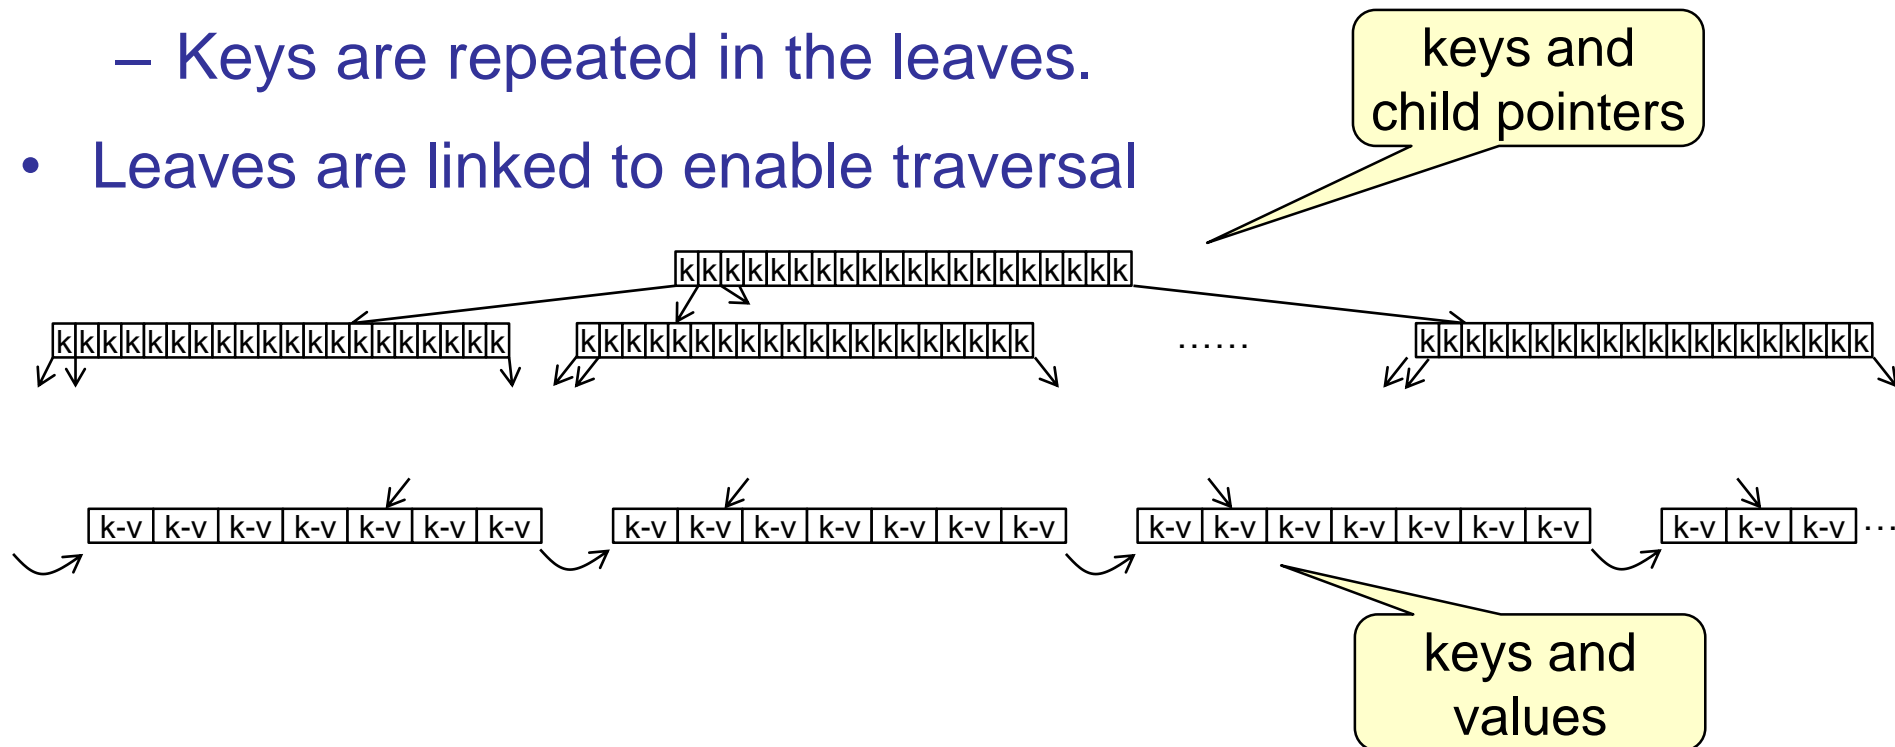
# Traversing a B Tree

- Listing items in order is expensive:
  - Moving up and down the tree means lots of record accesses.




# B+ Trees

- Commonly variant of B Trees – used in many DB/file systems.
- Intended for storing key–value (or key–record) pairs in files.
- Leaves contain key-value/record/index pairs.
- Internal nodes only contain (some) keys – only for searching.
  - Keys are repeated in the leaves.
- Leaves are linked to enable traversal



# B+ Trees: Leaves

- Each leaf node contains
  - between  $\lceil \max_L / 2 \rceil$  and  $\max_L$  key-value pairs,
  - a link to the next leaf in the tree




- Keys are ordered: for each key  $K_i$  in the leaf :
$$K_i < K_{i+1}$$
- The value might be either
  - the actual associated value/record (if small enough)
  - the index of a data block where value can be found (maybe in another file)

# B+ Trees: Internal Nodes


Except root may have fewer

- Each internal node has
  - between  $\lfloor \max_N / 2 \rfloor$  and  $\max_N$  keys, and
  - up to  $\max_N + 1$  child node indexes



- Branching factor =  $\max_N + 1$
- Keys act as separators for subtrees.
  - For each key  $X$  in the subtree at  $C_i$ :


$$K_i \leq X < K_{i+1}$$



- $K_i$  is the leftmost key in the subtree at  $C_i$   
ie, the first key in leftmost leaf of  $C_i$

# B+ example

- 3-degree, Add in order: M H T S R Q B A F D Z



# B+ Tree: Find


To find value associated with a key:

Find(key):


```
if root is empty return null
else return Find(key, root)
```

Find(key, node):

```
if node is a leaf
  for i from 0 to node.size-1
    if key = node.keys[ i ] return node.values[ i ]
  return null
if node is an internal node
  for i from 1 to node.size
    if key < node.keys[ i ] return Find(key, getNode(node.child[i-1]))
  return Find(key, getNode(child[node.size] ))
```



Could use  
binary search






# B+ Tree Add

---

- Find leaf node where item belongs
- Insert in leaf.
- If node too full,  
    split, and promote middle key up to parent,  
    middle key also goes to the right
- If root split, create new root containing promoted key

# Splitting a B<sup>+</sup>-Tree Leaf


- If a leaf overflows:
  - Leave the left most  $m$  keys in the node,
  - Move the right most  $m + 1$  keys to a new node,
  - Propagate the  $(m + 1)$ -st key to the parent node




- A non leaf node splits as in an ordinary B-tree
- The right sub-tree of each non leaf node contains greater or equal key values

# B<sup>+</sup>-Tree Insertion Example

- Given a B<sup>+</sup>-tree

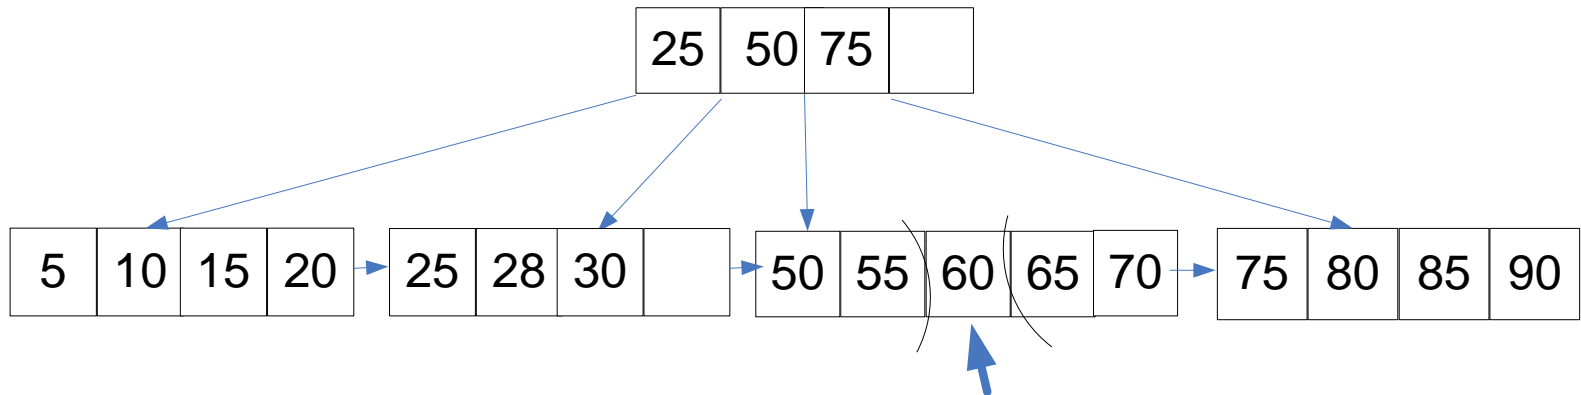



- Insert key 28




# B<sup>+</sup>-Tree Insertion Example (cont.)

- Insert key 70




# B<sup>+</sup>-Tree Insertion Example (cont.)

- The middle key of 60 is placed in a new node between 50 and 75

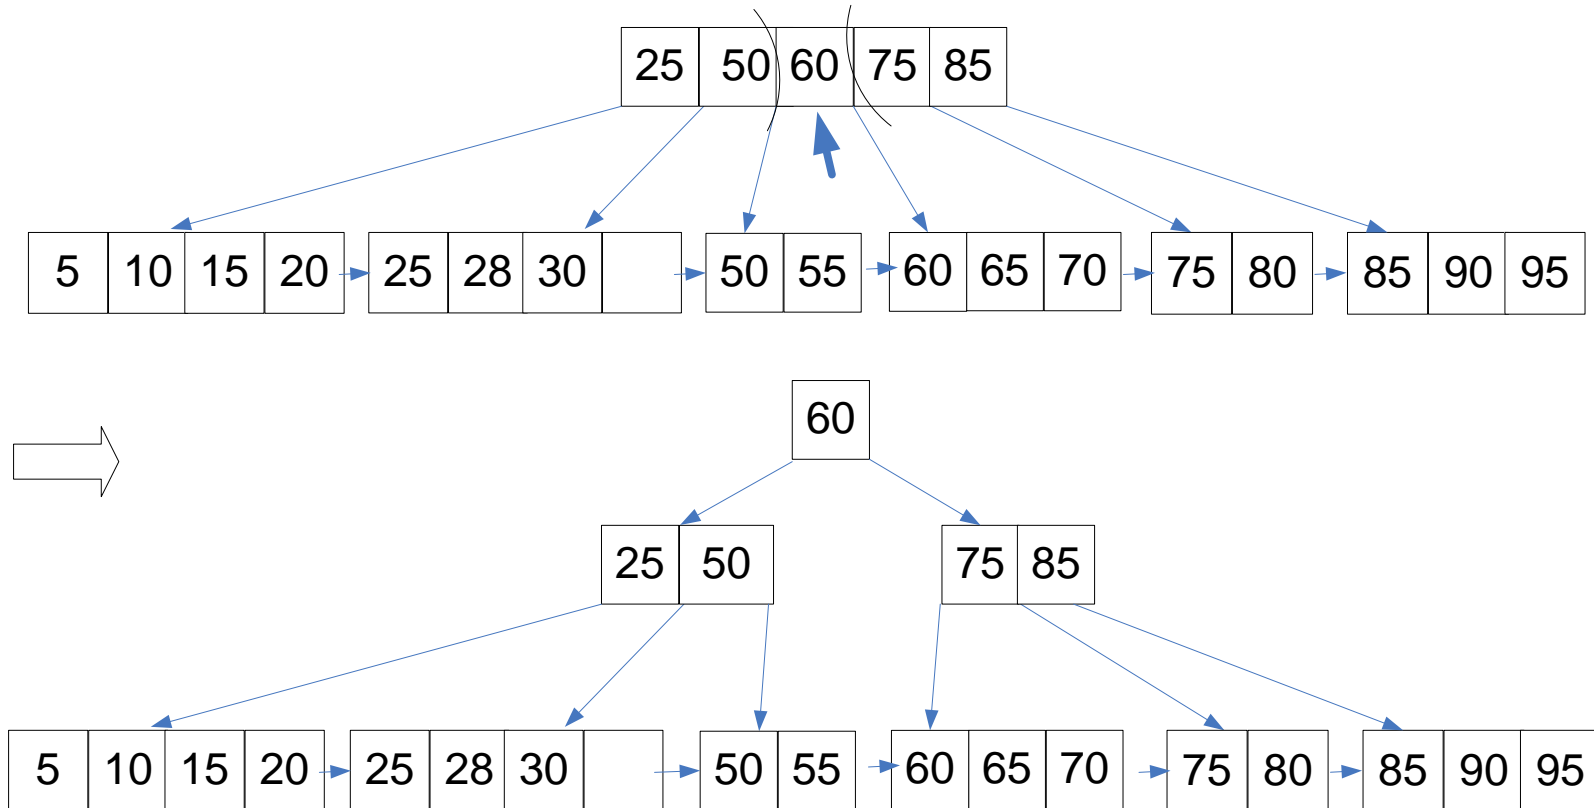


- Insert 95



# B<sup>+</sup>-Tree Insertion Example (cont.)

- Split the leaf and promote the middle key to the parent node



# B+ Tree Add (1)

Add(key, value):

**if** root is empty  
    create new leaf, add key-value,  
    root  $\leftarrow$  leaf

**else**

(newKey, rightChild)  $\leftarrow$  Add(key, value, root)

**if** (newKey, rightChild)  $\neq$  null

    node  $\leftarrow$  create new internal node  
    node.size  $\leftarrow$  1  
    node.child[0]  $\leftarrow$  root  
    node.keys[1]  $\leftarrow$  newKey  
    node.child[1]  $\leftarrow$  rightChild  
    root  $\leftarrow$  node

If root was full:  
returns new key  
and new leaf node,

Make a new  
root node

# B+ Tree Add (2)

Add(key, value, node):

**if** node is a leaf

**if** node.size < maxLeafKeys

        insert key and value into leaf in correct place

**return** null

**else**

**return** SplitLeaf(key, value, node)

$K_0-V_0, K_1-V_1, K_2-V_2, \dots, K_{\text{size}-1}-V_{\text{size}-1}$

Returns new key  
and new leaf node,

---

**if** node is an internal node

**for** i from 1 to node.size

**if** key < node.keys[i]

$(k, rc) \leftarrow \text{Add}(\text{key}, \text{value}, \text{node.child}[i-1])$

**if**  $(k, rc) = \text{null}$  **return** null

**else** **return** dealWithPromote(k, rc, node)

$(k, rc) \leftarrow \text{Add}(\text{key}, \text{value}, \text{node.child}[\text{node.size}])$

**if**  $(k, rc) = \text{null}$  **return** null

**else** **return** dealWithPromote(k, rc, node)

$C_0, K_1 C_1, K_2 C_2, \dots, K_{\text{size}} C_{\text{size}}$



Inserts new  
key and child  
into node,



# B+ Tree Add (3)

Could make the array one larger than necessary to give room for this.

SplitLeaf(key, value, node):

insert key and value into leaf in correct place (spilling over end)

sibling  $\leftarrow$  create new leaf

mid  $\leftarrow \lfloor (\text{node.size}+1)/2 \rfloor$

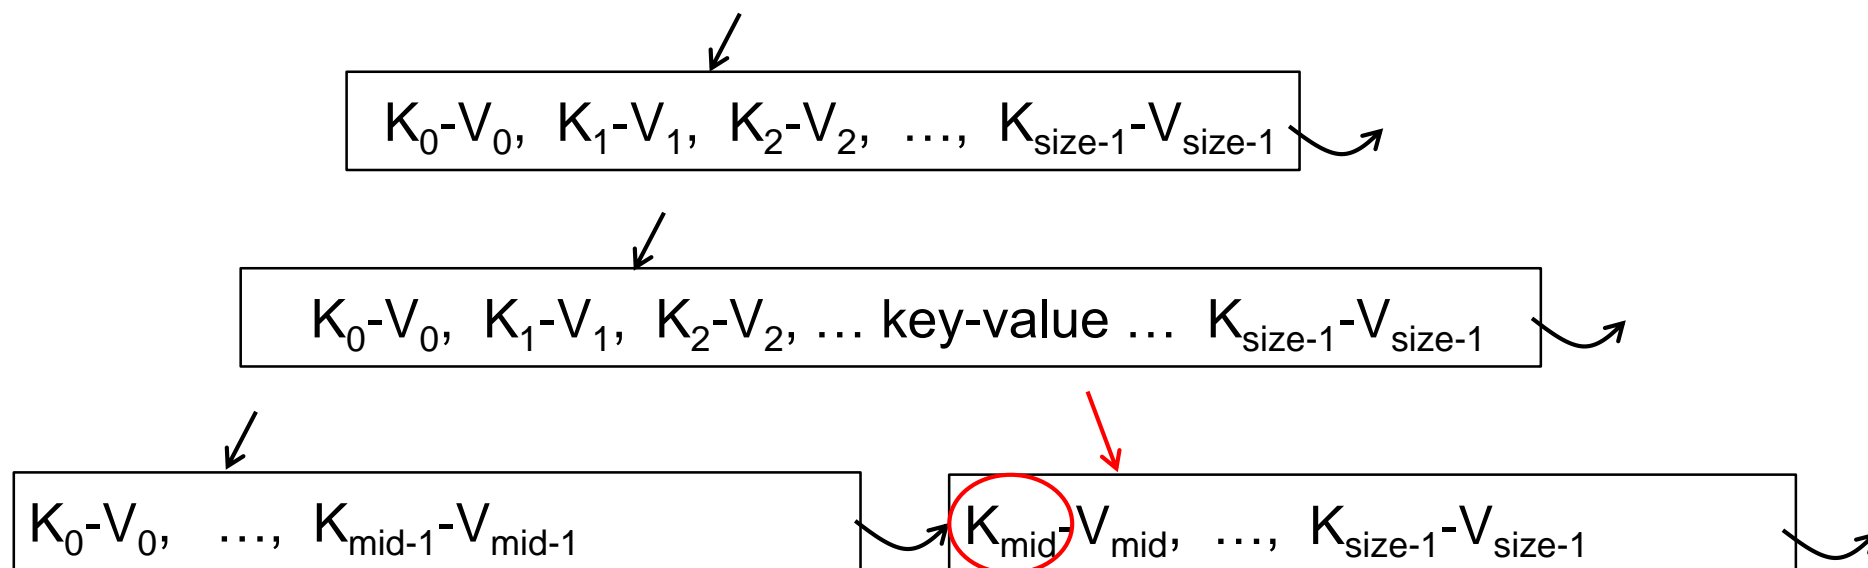
=  $\lfloor (\text{max}_L+2)/2 \rfloor$  since size is now  $\text{max}_L+1$

move keys and values from mid ... size out of node into sibling.

sibling.next  $\leftarrow$  node.next    node.next  $\leftarrow$  sibling



return (sibling.keys[0] , sibling)



# B+ Tree Add (4)

DealWithPromote( newKey, rightChild, node ):

**if** (newKey, rightChild) = null **return** null

**if** newKey > node.keys[node.size]

    insert newKey at node.keys[node.size+1]

    insert rightChild at node.child[node.size+1]

**else for** i from 1 to node.size

**if** newKey < node.keys[i]

        insert newKey at node.keys[i]

        insert rightChild at node.child[i]

**if** size  $\leq$  maxNodeKeys **return** null

sibling  $\leftarrow$  create new node

mid  $\leftarrow \lfloor \text{size}/2 \rfloor + 1$

move node.keys[mid+1... node.size] to sibling.node [1... node.size-mid]


move node.child[mid ... node.size] to sibling.child [0 ... node.size-mid]

promoteKey  $\leftarrow$  node.keys[mid],

remove node.keys[mid]

**return** (promoteKey, sibling)

Nothing was promoted



No need to promote further



Node is overfull:  
Have to split and promote

# B<sup>+</sup>-Tree Deletion

- When a record is deleted from a B<sup>+</sup>-tree it is always removed from the leaf level
- If the deletion of the record does not cause the leaf underflow
  - If the key of the deleted record appears in an index node, use the next key to replace it
- If deletion causes the leaf and the corresponding index node underflow
  - Redistribute, if there is a sibling with more than  $m$  keys
  - Merge, if there is no sibling with more than  $m$  keys
  - Adjust the index node to reflect the change



# B<sup>+</sup>-Tree Deletion Example

- Delete 70






# B<sup>+</sup>-Tree Deletion Example (cont.)

- Delete 25



# B<sup>+</sup>-Tree Deletion Example (cont.)

- Delete 60



# Using B+ trees for organising Data


- B+ tree is an efficient index for very large data sets
- The B+ tree must be stored on disk (ie, in a file)
  - ⇒ costly actions are accessing data from disk
- Retrieval action in the B+ tree is accessing a node
  - ⇒ want to make one node only require one block access
  - ⇒ want each node to be as big as possible ⇒ fill the block

B+ tree in a file:


- one node (internal or leaf) per block.
- links to other nodes = index of block in file.
- need some header info.
- need to store keys and values as bytes.

# Implementing B+ Tree in a File

- Use a block for each node of the tree
  - Can refer to blocks by their index in the file.
- Need an initial block (first block in file) with meta data:
  - index of the root block
  - number of items
  - information about the item sizes and types?




- Use a block for each internal node
  - type
  - number of items
  - child **key** child **key**.... **key** child



index of block  
containing child node

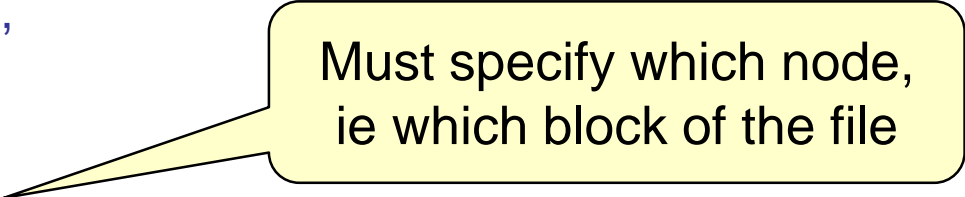
- Use a block for each leaf node
  - type
  - number of items
  - link to next
  - **key-value key-value .... key-value**



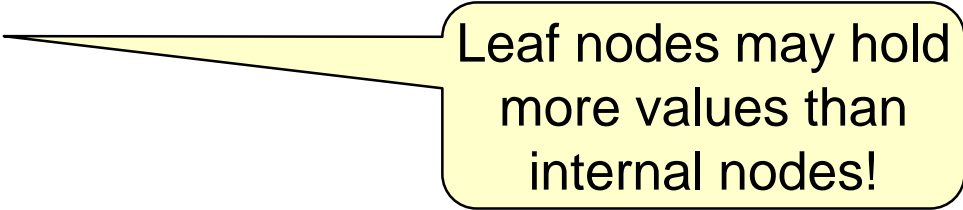


# Cost of B+ tree

- If the block size is 1024 bytes, how big can the nodes be?
- Node requires
  - some header information
    - leaf node or internal node
    - number of items in node,
  - internal node:
    - $m_N$  x key size
    - $m_N + 1$  x pointer size
  - leaf node
    - $m_L$  x item size
    - pointer to next leaf
- How big is an item?
- How big is a pointer?



Must specify which node,  
ie which block of the file



Leaf nodes may hold  
more values than  
internal nodes!

# Cost of B+ tree: Example

- Suppose:
  - a block has 1024 bytes
  - each node has header  $\Rightarrow$  5 bytes
  - a key is a string of up to 10 characters  $\Rightarrow$  10 bytes
  - a value is a string of up to 20 characters  $\Rightarrow$  20 bytes
  - a child pointer is an int  $\Rightarrow$  4 bytes
- Internal node ( $m_N$  keys,  $m_N+1$  child pointer)
  - size =  $5 + (10 + 4) m_N + 4$
  - $\Rightarrow m_N \leq (1024 - 9) / 14 = 72.5 \quad \Rightarrow 72 \text{ keys in internal nodes}$
- Leaf node (with pointer to next)
  - size =  $5 + 4 + (10 + 20) m_L$
  - $\Rightarrow m_L \leq (1024 - 9) / 30 = 33.8 \quad \Rightarrow 33 \text{ key-value pairs in leaves}$